

# ParSA: High-throughput Scientific-data Analysis Framework with Distributed File System

Tao Zhang<sup>a,\*</sup>, XiangZheng Sun<sup>b</sup>, Wei Xue<sup>a</sup>

<sup>a</sup>*Department of Computer Science and Technology, Tsinghua University, Beijing, China*

<sup>b</sup>*Intel Corporation, Beijing, China*

---

## Abstract

Scientific-data analysis and visualization has become a key component in nowadays large-scale simulations. Due to the rapidly increasing data volume and awkward I/O pattern among high-structured files, known serial methods/tools cannot scale well and usually lead to poor performance over traditional architectures. In this paper, we propose a new framework: ParSA (parallel scientific-data analysis) for high-throughput and scalable scientific analysis, with distributed file system. ParSA presents the optimization strategies for grouping and splitting logical units to maximize distributed I/O property of distributed file system as well as to maximize overlapping the data reading, processing and transferring during computation. Besides, ParSA provides the similar interfaces as the NCO (NetCDF Operator), which is used in most of climate data diagnostic packages, making it easy to port this framework. We utilize ParSA to accelerate well-known analysis methods for climate models on Hadoop Distributed File System (HDFS). Experimental results demonstrate the high efficiency and scalability of ParSA, getting the maximum 1.2GB/s throughput.

## Keywords:

Data intensive, Scientific data analysis, Distributed file system

---

---

\*Corresponding author. Tel: +86 18210192568

Email address: [t-zhang11@mails.tsinghua.edu.cn](mailto:t-zhang11@mails.tsinghua.edu.cn) (Tao Zhang)

## 1. Introduction

In most of modern scientific applications, huge amounts of data are produced. Large-scale simulations, such as climate modeling, high-energy physics simulation and genome mapping, generate hundreds of terabytes data volumes(Tevfik, 2009; Hey, 2003). Additionally, it still increases as the high resolution model developping. As a consequence, analysis of scientific-data is data-intensive.

In fact, almost all of scientific-data are stored in high-structured files, some of which provide parallel I/O interface, such as Network Common Data Format version 4(NetCDF4), Hierarchical Data Format 5(HDF5) (hdfs) and ADIOS BP data format (bp) (bp), and some of which only support serial I/O interface, like Network Common Data Format version 3(NetCDF3). All of these are self-describing, machine-independent data format.

In scientific-data analysis, large-scale scientific-data are stored in RAID-5/6 or parallel file system. Yet the analysis methods/ tools are always centralized approaches, such as NCO and NCL, which are the most used in climate applications for processing NetCDF files, leading to very poor scalability and time-consuming performance.

Inspired by big data solution in Internet Port Data Center (IPDC), numerous frameworks with distributed strategy have been developed. MapReduce is a program framework for processing and generating large data sets, provideing automatic parallel mechanism and build-in fault-tolerance on a cluster. However, this solution with MapReduce requires that data first be transformed into a text-based format. SciHadoop is a Hadoop plugin allowing scientists to specify logical queries over array-based data models. It executes queries as map/reduce programs defined over the logical data model. It shows remarkable improvements for holistic functions of NetCDF data sets for the following optimization goals: reduce total data transfers, reduce remote reads and reduce unnecessary reads. Nevertheless, SciHadoop using java language leads to the compatibility problem to the existing climate data analysis tools, which is written by c shell scripts, NCL and NCO commands. The SWAMP project [9] has provided the parallel NCO operations, but the reading performance is still bottleneck.

In this paper, we propose a new framework - Parallel Scientific-data Analysis (ParSA). We utilize the distributed I/O property with Hadoop Distributed File System (HDFS) to improved data reading throughput. What is more, ParSA optimizes the data layout schedule stored in the distributed

38 file system to overlap the data reading, processing and transferring. Besides,  
39 it provides parallel NCO operations, cooperating with HDFS, making it easy  
40 to use the efficient tool, without changing a lot for current climate analysis  
41 package.

## 42 **2. HDFS and scientific-data analysis**

43 In this section, we will present the property of distributed file system  
44 HDFS, replicas and scheduler, which can be taken advantage of to optimize  
45 distributed I/O performance. We also present the character of scientific-  
46 data analysis, and discuss about the probability of analysis transportaion  
47 onto HDFS.

### 48 *2.1. HDFS*

49 HDFS is an open source project, driven by Google File System (GFS).  
50 As a distributed, scalable and portable file system, HDFS is inherent for  
51 large-scale data-intensive process.

52 In HDFS, there are two types of node, Namenode and Datanode. Namen-  
53 ode maintains file system tree and metadata for all files or directories stored  
54 in HDFS, and Datanodes are where the data are actually stored. When a  
55 file are stored into HDFS, it will be split into file blocks as the storage unit  
56 of HDFS. For achieving fault-tolerance, HDFS stores three replicas, by de-  
57 fault, for each file blocks in different datanodes. Therefore, even individual  
58 node halts down, all data, which are stored in the halted node, can be ac-  
59 cessed from other replicas. All of I/O operates can be mantipulaed through  
60 Clientnode.

61 Each Datanode can mount several hard disks and it manages these hard  
62 disks by itself. By default, Datanode will store each block into the hard  
63 disks in a round-robin way. For example, when a file “a” will be stored into  
64 HDFS via Clientnode, Namenode will add “a” into the file system tree. Then  
65 Clientnode begin writing the content into HDFS. Once Clientnode detects  
66 that current writing size exceeds the block size, 64MB by default, it will  
67 ask Namenode for a new block with unique block number. Simultaneously,  
68 Namenode need recode mapping relation between “a” and block numbers.  
69 Since three replicas are used in HDFS for fault tolerance, Namenode will  
70 select three Datanodes to store a block in the file “a”. In Datanode, HDFS  
71 should choose a hard disk to each relic of a block. As shown in Figure 1,

the last block “k” is stored in disk “2”, then the new one “j” will locate at disk “0” according to the round-robin rule.

Blocks of files are distributed in the two-tiered storage architecture, Datanode and Datanode’s disk. It will make full use of the collective bandwidth of HDFS if each replica of blocks is appropriately chosen and scheduled to reduce remote reading among Datanodes.

MapReduce is usually introduced as a computation model cooperating with HDFS. It can process the data with good locality. However, due to the traits of scientific data and the operation pattern, MapReduce cannot be directly utilized on scientific-data analysis with perfect performance.

## *2.2. Scientific-data Analysis*

Scientific-data are usually stored in high-structured files, a kind of special binary format, which can not be read directly by MapReduce. Each file contains multi logical units, and each unit corresponds to its own physical meaning. The size of each logical unit varies a lot in one file and only some logical units can be manipulated in most of analysis operations.

In this paper, climate model data is used, the total size of which is extremely huge. The high resolution ocean model, a sub-component model in climate model, has 48TB for 100 simulated years, with monthly output. These data are organized by thousands of files with timestamp named in NetCDF format, called history result. Each file contains the same dozens of physical variables. Analysis operations need to process the whole or part of the history results. These operations include computing average, combining the same variables in different files, producing the remapping file and fast Fourier transform (FFT), etc. They require to handle whole or part of variables in multi-files.

## *2.3. Problems and Challenges*

Problem/Challenge 1: How to define the operation unit considering the storage unit. Logical unit size varies a lot from a few bytes to gigabytes. Then it can not be used as the operation unit. Otherwise, the workflow of each task will be imbalanced.

Problem/Challenge 2: How to schedule tasks among many files to increase the locality and reduce the data network transfer? If all logical units with same meanings locate on the same Datanode, all data access will proceed locally. However, default block placement policy of HDFS cannot ensure that, which leads to huge amount of data network transfer, as the above mentioned.

Problem/Challenge 3: How to best utilize the disk I/O?

### 3. Parallel scientific-data analysis

In order to solve the problems mentioned in previous section to take full use of distributed I/O performance, we propose a new framework - Parallel Scientific-data Analysis (ParSA).

#### 3.1. Logical Unit Split and Group among File Blocks

No matter how a file is split, all data in one block are stored contiguously in a physical location, since file block is the storage unit of HDFS. Thus the block of HDFS can be used as the basic operation unit to keep continuity in disk and improve I/O performance by proper scheduling to reduce network read. In a block, we can group the small size of logical units or split the big size one.

The logical unit split and group approach is shown in Figure 2. All logical units located in one block are distributed into one group. If a logical unit spans two or more blocks, this logical unit will split to several parts. For example, LU1 spans three blocks. It will split into three split parts - LU1 (1), LU1 (2) and LU1 (3), unique identifier assigned to each part. Then LU1 (1) and LU0 locate in file block 0 and form a operation unit group. LU1 (2) itself is distributed as a group. The remaining LU1 (3), LU2 and LU3 are distributed to the third group. The mapping relationship between groups and file blocks should be records, as shown in Figure 3.

Variables in NetCDF, used in climate data, can be viewed as multi-dimension array. An array section is defined as a contiguous rectangular block specified by two vectors - index vector and count vector. The index vector indicates the start offset of the element in the corner closest to the origin. The count vector gives the lengths of the edges along each of the variable's dimension.

For instance, the variable given in Figure 4 (a) is a two-dimension slab. If we access the whole variable, index vector is (0, 0) and count vector is (3, 5). However, if this variable locates on two file blocks stored with HDFS, and variable are split between value (1, 1) and value (1, 2), it cannot be split into these two pieces. The reason is that each of part is not a rectangular shape, which can not use an index vector and count vector to represent. In this situation, we should split the variable into four pieces, as Figure 4(b) shows.

### 143 3.2. Scheduler

144 As discussed in section 2.3, quite a few operations are only executed on  
 145 the same logical units among multi-files. In principle, blocks containing these  
 146 units should be scheduled in one Datanode as many as possible to increase  
 147 the access locality.

148 At first, information about file distribution should be gathered. As dis-  
 149 cussion in section 2.1, file block has three replicas by default. In this paper,  
 150 we define the tuple  $(N_i, D_j)$  as the location of a file block, which means  
 151 Datanode number and hard disk number, respectively. Assume we get the  
 152 following block location information for the first file block of all three files.

153 For the nine replicas for B0 of F1, F2 and F3, Datanode N0 covers 2  
 154 replicas. N1 contains 2 replicas and other Datanode only contain 1 for each.  
 155 Then Datanode N0 and N1 contain the same number of replicas. Therefore,  
 156 Datanode N0 is chosen to process the B0 block of F1 and F2, instead of N1  
 157 to process B0 of F0 and F2. The reason is that B0 of F0 and F2 locates  
 158 on the same hard disk in N1 Datanode, which is more likely to impact the  
 159 performance (explained in section 3.3). Although N1, N6 and N7 all have  
 160 the B0 replica of F2, N1 is selected. Because when disk 0 of N0 is broken,  
 161 the replica B0 of F0 can be accessed in N1. According to this principle, other  
 162 blocks can be scheduled as the Figure 5(b) shown.

### 163 3.3. Workflow of ParSA

164 Although we utilize the file block location in HDFS to schedule which  
 165 block replica are selected among Datanodes, I/O bandwidth of disks may not  
 166 be efficiently utilized intra-Datanode. For example, the file block B0s of F0  
 167 and F1 are assigned to N0 in Figure 5. Even though B0 of F0 and B0 of F1  
 168 locate on different disks - D0 and D1 respectively, the two disks are accessed  
 169 sequentially instead of parallelly, if the operations are implemented from logical  
 170 level at Datanode N0 as following:

---

#### **Algorithm 1** Sequential operation on disk in the same Datanode

---

```

for each logical-unit LU in B0 do
  for each file Fx whose block B0 locates on N0 do
    read LU from B0 of Fx
    operate on LU
  end for
end for

```

---

171 In this paper, we propose a new workflow to efficiently access the data  
 172 on HDFS and maximize utilizing all disk bandwidth in one Datanode.

173 In ParSA, the framework of hybrid of MPI and pthread is used. MPI is  
 174 responsible for parallel and communication among Datanodes. There is one  
 175 MPI process in each Datanode. Pthead is responsible to implement the detail  
 176 operations within each Datanode, constituting of three major components:  
 177 reading-thread, processing-thread and receiving-thread.

178 **Reading-thread** If more than one thread dedicatedly access one disk,  
 179 it cannot get a higher performance than only one thread. Thus we invoke  
 180 one thread for each disk. This is defined as reading-thread and use  $Tread_i$   
 181 to present the  $i_{th}$  reading-thread. For each block,  $Tread_i$  reads the queue  
 182 of logical units, shown in Figure 3. After reading the logical unit, it is  
 183 marked as LOCAL\_FLAG, distinguishing with receiving data. The workflow  
 184 of reading-thread is given below:

---

**Algorithm 2** reading-thread

---

inputs:

**datanode\_number**:the Datanode where the reading-thread locates.

**disk\_number**:the disk number to be accessed.

**scheduler\_result**:task scheduler result, given in Figure 5(b).

**LU\_groups\_map**:mapping relations between logical unit groups and file  
 blocks in Figure 3.

```

for each tuple (node_no, disk_no) in scheduler_result do
  if node_no = datanode_number and disk_no = disk_number then
    get the file-index and block-index according to the index of tuple,
    shown in Figure 5(a).
    for each lu in LU_groups_map[block_index] do
      read values of lu of file_index file
      store file-index and values into lu
      mark lu as LOCAL_TAG
      insert (lu) into data queue
      notify processing-thread
    end for
  end if
end for
send ending-flag to processing-thread

```

---

185     **Processing-thread** After reading-threads read logical units into queue,  
186 the processing-thread will be waked up. The processing- thread takes in  
187 charge of calling user-defined function `local_data_process` to process these  
188 logical units. Users only focus on implementing operations on logical unit.  
189 Since each Datanode only covers part of entire dataset, the major process  
190 should reduce the intermediate results, computed with each process in each  
191 Datanode. This operation is another user-defined `recv_data_process` function.

192     **Receiving-thread** Receiving-thread is a component in major process. It  
193 receives logical inits from worker processes, and puts the logical units into the  
194 receiving queue. Then the processing-thread is notified to fetch these data.

195     The workflow of ParSA is shown in Figure 6. Major process locates in  
196 Datanode 0, and a worker process is presented as Datanode2. Each read-  
197 ing thread reads the logical units which is assigned to this thread. After  
198  $Tread_0$  in Datanode 2 reading the logical LU1 (3) of F0, it reads LU2 of  
199 F0, inserts the values into queue and notifies processing-thread to wake up.  
200 The processing-thread processes all the values in queue sequentially. When  
201 processing-thread detects that LU1 (3) are all processed in Datanode 2, it will  
202 send the intermediate results to the major process. Once Receiving-thread  
203 in major process gets the intermediate results, it inserts the results into the  
204 logical unit queue. Since the logical unit queue is the shared resource, all  
205 threads must access exclusively.

206     ParSA is the framework for scientific-data access on HDFS. Users can  
207 only put their focus on logical operations. ParSA is implemented using MPI  
208 and Pthread, providing more flexible configuration with regard to performance  
209 than MapReduce. ParSA can achieve load-balance based on current block  
210 replicas layout information. However, if the layout itself is unbalance, the  
211 performance is also impacted. In the next section, how to further improve  
212 the performance by optimizing layout will be discussed.

## 213 4. Block layout optimization

214     In section 3.2, the scheduler of ParSA utilizes the location information of  
215 file blocks. In principle, file blocks containing the same logical units in the  
216 one Datanode should be scheduled as many as possible for the operations  
217 involving multi-files, while the default block layout strategy of HDFS cannot  
218 guarantee this.

219     In HDFS, if Clientnode is one of the Datanodes, the first replica of each  
220 file block will be stored in Clientnode, and other replicas are chosen random



221 Datanodes expect for Clientnode. If the Clientnode is not Datanode, the  
222 first replica will be randomly stored in a Datanode.

223 Although the scheduler can assign tasks based on current file block repli-  
224 cas distribution, the distribution should be control according to the principle  
225 for task assignment in order to achieve load balancing. Since blocks of files  
226 are stored in Datanode and Datanode's disk, it needs to optimize the layout  
227 in these the two tiered architecture.

#### 228 4.1. Data layout inter-Datanodes

229 Taken into account the characteristic of scientific-data analysis, tightly  
230 coupled data as above described should be assigned in one Datanode to  
231 increase locality of processing data. Simultaneously, each Datanode should  
232 have equal workload, with a pretty load-balance performance.

233 However, it may indeed exists a few analysis operations different from the  
234 characteristic. ParSA expose an interface for users to determine where the  
235 first replicas of blocks shoule locate. What's more, the second replicas are  
236 stored with uniform distribution automaticlly. The third abide by the HDFS  
237 default strategy.

#### 238 4.2. Data layout intra-Datanodes

239 As discussed in section 3.3, the workflow of ParSA invokes one thread  
240 for one disk. If blocks are imbalance among disks in one Datanode, it will  
241 waste the bandwidth of other disks. The load balance is not only the equal  
242 number of file blocks in one Datanode, but also the distribution of blocks  
243 containing the same logical units. The former can be guaranteed by HDFS  
244 default disk allocation strategy with a round-robin way, while the later can  
245 not be satisfied all the time. ParSA provides a round-robin mechanism for  
246 these blocks distribution, as shown in Figure 7. The disk will be used for a  
247 block, if it does not store one with the same logical units.

## 248 5. Evaluation

### 249 5.1. Experimental setting

250 To demonstrate the performance portability of ParSA framework, we  
251 choose two HDFS clusters with different configuration, as shown in Table  
252 1 and Table 2.

253 In order to demonstrate the efficiency of ParSA, we use climate data  
254 generated by a ocean model, POP, a component of climate model. The

255 dataset has 100 NetCDF files, and each of file contains 60 two-dimension  
256 variables (logical units) and 18 three-dimension variables, with 320x384 and  
257 320x384x40 resolution, which has 309 MB datasize. The total size of this  
258 dataset is 30.9 GB.

259 The major analysis operation is calculating time average using ncea com-  
260 mand of NCO package, which is the basic preprocessing operation for almost  
261 all of climate data analysis.

262 To evaluate the scalability of ParSA, we compare the speedup and through-  
263 put performance of different number of Datanodes and different number of  
264 disks per node. Cluster 1 is used for the former experiment since the number  
265 nodes of cluster 1 is more than cluster 2. While cluser2 is used for the later  
266 experiment since each node has more disks in cluster 2.

### 267 5.2. Scalability with different disks per Datanode

268 ParSA can get about 1.3 GB/s aggregation throuphout if per Datan-  
269 ode uses 5 disks, as shown Figure 8(a). The throughput has a remarkable  
270 scalability if per Datanode from 1 to 4 disks. However, it has little perfor-  
271 mance improvement with 5 disks. The reason is that there are some serial  
272 operations, such as writing results and communication, which impact on the  
273 performance. Since a series of strategies propused by ParSA is aimed at I/O,  
274 Figure 8(b) shows the speedup with only reading part. It's scalability is close  
275 to the ideal speedup.

### 276 5.3. Scalability with different Datanodes

277 Two disks per node are used in this evaluation with cluster 1. It almost  
278 keeps the linear speedup until 9 Datanodes as shown Figure 9. However, the  
279 performance improvment is limited when the datanodes are greater than 9  
280 due to the increasing overhead of communication and the critical section con-  
281 trolled by mutex, as shown Figure 6. Although data processing is overlapped  
282 with data accessing, the time non-overlapped occupies almost 38.89% on 16  
283 datanodes, and this proportions are only 0.02% and 7.96% when datanodes  
284 are 9 and 12 respectively.

285 In Figure 9(a), it can get about 800MB/s maximum throughput on 12  
286 datanodes, with 24 disks totally. While the same number of disks in Figure  
287 8(a), 4 disks per datanode, it can get about 1.2GB/s. The same performance  
288 only need 2 disks per datanode, with 12 disks totally. The reason is that, on  
289 one hand, the disks and network are difference between the two cluters; on

290 the other hand, cluster 1 requires more network access due to the few disks  
291 per node.

292 The speedup of only reading is presented as Figure 9(b), which shows an  
293 excellent scalability with different datanodes. The reason why the speedup is  
294 higher than the ideal may result from performance variability.

#### 295 5.4. *Throughput comparison with different task optimization strategies*

296 HDFS block is the basic operation unit in ParSA, splitting and grouping  
297 the logical unit. We select the block replicas in different datanodes to reduce  
298 network reading, optimize the workflow to support overlapping and paral-  
299 lelism of reading and processing, as well as the optimization of blocks layout  
300 for both inter-datanodes and intra-datanodes.

301 To present the improving performance, two additional strategies are com-  
302 pared to ParSA. The first one takes logical unit as the operation unit. The  
303 disk utilization is shown as Figure 10(a). Since the size of logical unit varies  
304 quite a lot, a serious load imbalance exists among tasks. The second strat-  
305 egy takes all logical units located in one file block as an operation unit, while  
306 those logical unit which locate on more than one block do not be split and  
307 treat the big size logical units as other operation unit. In this situation, it  
308 can not guarantee that an entire operation unit stores in a disk. When a  
309 thread try to read this unit, it requires read a part of value from the other  
310 disk. In the mean time, the other thread is reading the same disk, leading  
311 to disk resource competition. This utilization performance is presented as  
312 Figure 10(b).

313 Compared with throughput to evaluate the speedup, ParSA outperforms  
314 the above two strategies by a factor of 3.34X and 1.14X respectively.

## 315 6. Conclusion and future work

316 A high-throughput and scalable scientific framework combined with HDFS  
317 distributed file system and NCO analysis tool has been designed, imple-  
318 mented and evaluated. Maximize distributed I/O performance has been  
319 gained by scheduling block replicas in different datanodes, optimizing the  
320 blocks layout for both inter-datanodes and intra-datanodes, as well as over-  
321 lapping the data reading, processing and transferring. Instead of Hadoop  
322 framework, we use multi-process with MPI to parallel the NCO basic oper-  
323 ations, without leading to compatible problem. In future work, we plan on

324 testing the framework on some complex climate data analysis, such as FFT,  
325 principal component analysis and Wavelet transform.

326 [1]

327 [2]

---

**Algorithm 3** processing-thread

---

inputs:

**num\_tasks**:total number of **ending-flag** to wait

**intermediate\_results**:store the intermediate results

**cur\_num\_tasks** = 0

**while** **cur\_num\_tasks** < **num\_tasks** **do**

    wait until receiving notification

**if** receive **ending-flag** **then**

**cur\_num\_tasks**++

        fetch the logical unit queue

**for** each **lu** in logical unit queue **do**

            get **interm\_res** according to **lu identifier** from **intermediate\_results**

**if** **lu's flag** == **LOCAL\_FLAG** **then**

                call **local\_data\_process** (**lu**, **interm\_res**)

**end if**

**if** **lu's flag** == **RECV\_FLAG** **then**

                call **recv\_data\_process** (**lu**, **interm\_res**)

**end if**

**if** is not major process **then**

**if** all logical unit are processed in this Datanode **then**

                    store **interm\_res** into **lu**

                    mark **lu** as **RECV\_TAG**

                    send (**lu**) to major process

**end if**

**end if**

**end for**

**end if**

**end while**

---