

# ParSA: High-throughput Scientific-data Analysis Framework with Distributed File System

Tao Zhang

*Tsinghua*

---

## Abstract

Scientific-data analysis and visualization has become a key component in nowadays large-scale simulations. Due to the rapidly increasing data volume and awkward I/O pattern among high-structured files, known serial methods/tools cannot scale well and usually lead to poor performance over traditional architectures. In this paper, we propose a new framework: ParSA (parallel scientific-data analysis) for high-throughput and scalable scientific analysis, with distributed file system. ParSA present the optimization strategies oriented for physical disk to maximize distributed I/O property of distributed file system as well as to maximize overlapping the data reading, processing and transferring during computation. Besides, ParSA provides the similar interfaces as the NCO (netCDF Operator), which is used in most of climate data diagnostic package, making it easy to port this framework. We utilize ParSA to accelerate well-known analysis methods for climate models on Hadoop Distributed File System (HDFS). Experimental results demonstrate the high efficiency and scalability of ParSA.

*Keywords:*

Data intensive, Scientific data analysis, Distributed file system

---

## 1. Introduction

In most of modern scitific applications, huge amounts of data are produced. Large-scale simulations, such as climate modeling, high-energy physics simulation and genome mapping, generate hundreds of terabytes data volumes(Tevfik, 2009; Hey, 2003). Additionally, it still increases as the high resolution model developping. As a consequence, analysis of scientific-data is data-intensive.

8 In fact, almost all of scientific-data are stored in high-structured files,  
9 some of which provide parallel I/O interface, such as Network Common Data  
10 Format version 4(NetCDF4), Hierarchical Data Format 5(HDF5) (hdfs) and  
11 ADIOS BP data format (bp) (bp), and some of which only support serial  
12 I/O interface, like Network Common Data Format version 3(NetCDF3). All  
13 of these are self-describing, machine-independent data format.

14 In scientific-data analysis, large-scale scientific-data are stored in RAID-  
15 5/6 or parallel file system. Yet the analysis methods/ tools are always central-  
16 ized approaches, such as NCO and NCL, which are the most used in climate  
17 applications for processing NetCDF files, leading to very poor scalability and  
18 time-consuming performance.

19 Inspired by big data solution in Internet Port Data Center (IPDC), nu-  
20 merous frameworks with distributed strategy have been developed. MapRe-  
21 duce is a program framework for processing and generating large data sets,  
22 provideing automatic parallel mechanism and build-in fault-tolerance on a  
23 cluster. However, this solution with MapReduce requires that data first be  
24 transformed into a text-based format. SciHadoop is a Hadoop plugin al-  
25 lowing scientists to specify logical queries over array-based data models. It  
26 executes queries as map/reduce programs defined over the logical data model.  
27 It shows remarkable improvements for holistic functions of NetCDF data sets  
28 for the following optimization goals: reduce total data transfers, reduce re-  
29 mote reads and reduce unnecessary reads. Nevertheless, SciHadoop using  
30 java language leads to the compatibility problem to the existing climate data  
31 analysis tools, which is written by c shell scripts, NCL and NCO commands.  
32 The SWAMP project [9] has provided the parallel NCO operations, but the  
33 reading performance is still bottleneck.

34 In this paper, we propose a new framework - Parallel Scientific-data Anal-  
35 ysis (ParSA). We utilize the distributed I/O property with Hadoop Dis-  
36 tributed File System (HDFS) to improved data reading throughput. What  
37 is more, ParSA optimizes the data layout schedule stored in the distributed  
38 file system to overlap the data reading, processing and transferring. Besides,  
39 it provides parallel NCO operations, cooperating with HDFS, making it easy  
40 to use the efficient tool, without changing a lot for current climate analysis  
41 package.

## 2. HDFS and scientific-data analysis

In this section, we will present the property of distributed file system HDFS, replicas and scheduler, which can be taken advantage of to optimize distributed I/O performance. We also present the character of scientific-data analysis, and discuss about the probability of analysis transportation onto HDFS.

### 2.1. HDFS

HDFS is an open source project, driven by Google File System (GFS). As a distributed, scalable and portable file system, HDFS is inherent for large-scale data-intensive process.

In HDFS, there are two types of node, Namenode and Datanode. Namenode maintains file system tree and metadata for all files or directories stored in HDFS, and Datanodes are where the data are actually stored. When a file are stored into HDFS, it will be split into file blocks as the storage unit of HDFS. For achieving fault-tolerance, HDFS stores three replicas, by default, for each file blocks in different datanodes. Therefore, even individual node halts down, all data, which are stored in the halted node, can be accessed from other replicas. All of I/O operates can be manipulated through Clientnode.

Each Datanode can mount several hard disks and it manages these hard disks by itself. By default, Datanode will store each block into the hard disks in a round-robin way. For example, when a file “a” will be stored into HDFS via Clientnode, Namenode will add “a” into the file system tree. Then Clientnode begin writing the content into HDFS. Once Clientnode detects that current writing size exceeds the block size, 64MB by default, it will ask Namenode for a new block with unique block number. Simultaneously, Namenode need recode mapping relation between “a” and block numbers. Since three replicas are used in HDFS for fault tolerance, Namenode will select three Datanodes to store a block in the file “a”. In Datanode, HDFS should choose a hard disk to each replica of a block. As shown in Figure 1, the last block “k” is stored in disk “2”, then the new one “j” will locate at disk “0” according to the round-robin rule.

Blocks of files are distributed in the two-tiered storage architecture, Datanode and Datanode’s disk. It will make full use of the collective bandwidth of HDFS if each replica of blocks is appropriately chosen and scheduled to reduce remote reading among Datanodes.

MapReduce is usually introduced as a computation model cooperating with HDFS. It can process the data with good locality. However, due to the traits of scientific data and the operation pattern, MapReduce cannot be directly utilized on scientific-data analysis with perfect performance.

## 2.2. Scientific-data Analysis

Scientific-data are usually stored in high-structured files, a kinds of special binary format, which can not be read directly by MapReduce. Each file contain multi logical units, and each unit are corresponding to its own physical meaning. Since the size of each logical varies a lot in one file and only same logical units can be manipulated in some analysis operations, it can cause loading imbalance among works or too much network reading with MapReduce.

In this paper, climate model data is used, the total size of which is extremely huge. The high resolution ocean model, a sub-component model in climate model, has 48TB for 100 simulated years, with monthly output. These data are organized by thousands of files with timestamp named in NetCDF format, called history result. Each file contains the same dozens of physical variables. Analysis operations need to process the whole or part of the history results. These operations include computing average, combining the same variables in different files, producing the remapping file and fast Fourier transform (FFT), etc. They require to handle whole or part of variables in multi-files.

## 2.3. Problems and Challenges

Problem/Challenge 1: How to define the operation unit considering the storage unit - file block on HDFS? Logical unit size varies a lot from a few bytes to gigabytes, which means that one file block may contain hundreds of logical units whereas another file block only store small proportion of one logical unit. Then the logical block can not be used as the operation unit. Otherwise, the workflow of each task will be imbalance if each block contains the same logical units. Besides, it will increase much network reading if each block contains different logical units since most of operations in climate are manipulated with the same logical units in multi-files.

Problem/Challenge 2: How to schedule tasks among many files to increase the locality and reduce the data network transfer? If all logical units

112 with same meanings locate on the same Datanode, all data access will pro-  
113 ceed locally. However, default block placement policy of HDFS cannot ensure  
114 that, which leads huge amount data network transfer.

115 Problem/Challenge 3: How to best utilize the disk I/O?

### 116 **3. Parallel scientific-data analysis**

117 In order to solve the problems mentioned in previous section to take full  
118 use of distributed I/O performance, we propose a new framework - Parallel  
119 Scientific-data Analysis (ParSA).

#### 120 *3.1. Logical Unit Split and Group among File Blocks*

121 [1]